

Coursework - 4 : Report

Aim

To simulate a concurrent “banking” system. The main program should spawn ten “customer” threads, and each of these threads model a bank account with a starting balance of £1000. The customers should then (at random intervals) choose one of the other customers (at random) and transfer a random amount of money (between £10 and £50) into their account.

Executing the Program

Run the `stack build` to build the project and then `stack exec bank-exe` to execute it. Read more about that here in this project's [Github readme](#).

Control Flow and Explanation

- The program starts with starting the worker threads to create random transactions - which is done via forking the method to create transactions in a thread. We create 10 threads(each thread for each customer) and perform 100 transactions per thread. Mainly three things happen in a single forking(or in a single thread)
 - The data for transactions are being created randomly - two accounts to be debited and credited from (say A-> B) and the amount to be credited. Appropriate checks are added to that too.
 - The accounts are checked whether it exists, if not we create one with a minimum balance of £1000.
 - Now the transfer process happens, the amount is debited from A and credited to B. here it checks whether A!=B and after debiting A doesn't have a *negative balance*. If everything is fine, the transfer is initiated *atomically*.
- Now the main program waits for all the workers to finish their jobs
- And the amounts are mapped and printed in readable form.

P.S - The initial sum of all accounts is £10,000. So after these transactions the sums must be the same, ie £10,000. This can be used to verify whether the transactions are correct or not.

Design Choices and Issues faced

- We have used both `TVar` and `MVar` to avoid a few *deadlocking issues* which were caused in the initial stages of development. Source [post](#).
- We tried implementing the transactions more efficiently and real like by introducing the atomicity principle in Databases. It basically follows an all or nothing approach. It was implemented using the `STM` module. Source [link](#).
- Since we have used `TVar Mapping (Map.Map)` to create the custom datatype `Customer`, we had to use other mapping functions too like `Map.assocs`, `mapM`, `Map.empty` etc. Source posts - [one](#) and [two](#).
- *Issue (pending)* - The parallel mapping was tried to be implemented using `parallel_` and `parMap` to be used instead of replicating the `forkIO`. But it was creating issues in the existing computations and even running into locks (and even race conditions). Source [post](#).